

Task 2 Deliverable — Model Architecture (MRIEmbeddingModel)

Overview

This document describes the Task 2 model that maps batches of multi-coil complex MRI inputs into fixed-size embeddings. The implementation emphasizes simplicity, readability, and production-friendliness and is designed to work directly with Task 3 distance-based objectives.

I/O Contract

Input: $x \in \mathbb{C}^{(B \times C_{\text{in}} \times H \times W)}$, where $C_{\text{in}} = 2 \times \text{num_coils}$ (real and imaginary channels stacked by Task 1). Output: $y \in \mathbb{C}^{(B \times D)}$, where $D = \text{embed_dim}$ (L2-normalized by default).

Architecture Summary

- Stem: LazyConv2d \rightarrow GroupNorm \rightarrow GELU (adapts automatically to C_{in}).
- Backbone (4 stages): Conv2d(stride=2) \rightarrow GroupNorm \rightarrow GELU \rightarrow (optional Dropout).
- Head: GlobalAveragePool \rightarrow Linear(projection to embed_dim) \rightarrow (optional 2-layer MLP).
- Output: optional L2 normalization (enabled by default).

Key Design Choices & Rationale

- Compact CNN: satisfies the brief's emphasis on a simple, not overly complicated model while being fast and stable.
- GroupNorm: robust to small batch sizes common in medical imaging pipelines.
- LazyConv2d: avoids hardcoding the coil count; works with any C_{in} produced by Task 1.
- Global Average Pooling: parameter-efficient spatial aggregation.
- L2-normalized embeddings: distances (cosine/Euclidean) behave well for contrastive objectives in Task 3.
- Deterministic default: dropout is 0.0 by default to keep tests reproducible; enable during training if desired.

Initialization & Implementation Details

Convolutional layers use Kaiming initialization with ReLU gain (compatible with GELU). Linear layers use Xavier uniform. For the first LazyConv2d, parameters are initialized only after the first forward pass; the code guards against initializing unmaterialized lazy parameters.

Public API

File: model.py; Class: MRIEmbeddingModel(embed_dim=256, widths=(32,64,128,256), dropout=0.0, normalize=True, use_mlp_head=False). Methods: forward(x[, normalize]) \rightarrow (B, D); embed(x, normalize=True) \rightarrow (B, D) under no_grad.

Sanity Checks (PyTest)

A small test suite validates output shapes, L2 normalization behavior, gradient flow, and LazyConv2d adaptation to different channel counts. Command: PYTHONPATH=.

./venv/bin/python -m pytest -q. All tests pass.

Evaluation Criteria Compliance

- Accepts a batch of coil images and outputs a fixed-size embedding vector.
- Clean, readable, and production-friendly code with clear I/O contract and documentation.
- Simple architecture (compact CNN) that is easy to extend (e.g., optional MLP head).
- Rationale for design choices is documented for the deliverable write-up.

Usage Example

```
>>> import torch >>> from model import MRIEmbeddingModel >>> x = torch.randn(2, 16, 128, 128) # 8 coils → 16 channels (real+imag) >>> model = MRIEmbeddingModel(embed_dim=256, normalize=True) >>> y = model(x) # y.shape == (2, 256)
```

End of Document