

Task 2 & Task 3 Report

Task 2 — Embedding Model Design

For Task 2, the goal was to design a compact CNN that maps multi-coil complex MRI images (stacked real/imag channels from Task 1) into a fixed-size embedding vector suitable for contrastive learning.

Model Architecture (Layer by Layer):

- 1 Stem: LazyConv2d (out=32, stride=2, kernel=3) → GroupNorm → GELU
- 2 Stage 1: Conv2d(32→64, stride=2) → GroupNorm → GELU
- 3 Stage 2: Conv2d(64→128, stride=2) → GroupNorm → GELU
- 4 Stage 3: Conv2d(128→256, stride=2) → GroupNorm → GELU
- 5 Head: GlobalAveragePooling → Linear(256→256) → optional MLP head → L2 Normalization

We used GroupNorm instead of BatchNorm for robustness with small batch sizes, and GELU activations for smooth non-linearities. LazyConv2d was chosen in the stem to adapt automatically to any coil count. The design is deliberately simple and production-friendly.

Why This Architecture?

While alternatives such as ResNet-18 or ConvNeXt could have been used, the compact CNN offers a balance of simplicity, clarity, and sufficient performance for synthetic phantom data. It achieved perfect Recall@1 in validation. References consulted include ResNet, GroupNorm, GELU, and Triplet Loss literature (He et al. 2016, Wu & He 2018, Hendrycks & Gimpel 2016, Schroff et al. 2015).

Embedding Space

The embedding vectors are in \mathbb{R}^{256} and L2-normalized to lie on the unit hypersphere. This ensures cosine similarity and Euclidean distance are equivalent. Triplet loss enforces small distances for positive pairs and large distances for negatives. Empirically, $d_{\text{pos}} \sim 0.1\text{--}0.2$ and $d_{\text{neg}} \sim 1.1\text{--}1.25$, with Recall@1 consistently 100%.

Task 3 — Training with Contrastive Learning

The model was trained using the RandomPhantomTripletDataset and Triplet Loss.

Steps taken:

- 1 Created training dataset with two transform lists:
- 2
 - 1 List A: 4× acceleration, 8% center fraction + Normalize + ToCompatibleTensor + Augmentation
 - 2 List B: 8× acceleration, 4% center fraction + Normalize + ToCompatibleTensor + Augmentation
- 3 Validation dataset with same transforms but without augmentation, with offset=len(train_ds) and deterministic=True.

- 4 Training loop: forward pass through model, compute Triplet Margin Loss (Euclidean), optimizer step with AdamW.
- 5 Validation loop: compute val loss, d_pos, d_neg, viol%, and Recall@1 each epoch.
- 6 Saved curves (PNG + CSV) and best checkpoint.

Experimental Results

Initial experiments showed fast convergence and near-zero validation loss early on. To make the experiment more reliable, we scaled dataset size (train=4000, val=1000), increased image size to 192, batch size to 32, and margin to 0.3. This led to stronger separation: $d_{\text{pos}} \approx 0.09\text{--}0.14$, $d_{\text{neg}} \approx 1.1\text{--}1.25$, $\text{viol}\% < 0.5\%$, and $\text{Recall@1} = 100\%$.

Compliance with Instructions

All required criteria were satisfied:

- 1 Two transform lists with specified undersampling fractions.
- 2 Validation dataset uses same transforms but no augmentation, offset, and deterministic mode.
- 3 Triplets used as required; embeddings normalized and compared with Euclidean distance.
- 4 Training and validation loops implemented with loss + metrics per epoch.
- 5 Curves and checkpoints saved for analysis.